

Polynomial Curve Fitting


Polynomial Curve Fitting

A Simple Regression Problem


- We observe a real-valued input variable x and we wish to use this observation to predict the value of a real-valued target variable t .
- We use synthetically generated data from the function $\sin(2\pi x)$ with random noise included in the target values.
 - A small level of random noise having a Gaussian distribution
- We have a training set comprising N observations of x , written $\mathbf{x} \equiv (x_1, \dots, x_N)^T$, together with corresponding observations of the values of t , denoted $\mathbf{t} \equiv (t_1, \dots, t_N)^T$.
- Our goal is to predict the value of t for some new value of x ,

Polynomial Curve Fitting

- A training data set of $N = 10$ points, (blue circles),
- The green curve shows the actual function $\sin(2\pi x)$ used to generate the data.
- Our goal is to predict the value of t for some new value of x , without knowledge of the green curve.



Polynomial Curve Fitting




- We try to fit the data using a polynomial function of the form

$$y(x, \mathbf{w}) = w_0 + w_1x + w_2x^2 + \dots + w_Mx^M = \sum_{j=0}^M w_jx^j$$

Polynomial Curve Fitting

- The values of the coefficients will be determined by fitting the polynomial to the training data.
- This can be done by minimizing an error function that measures the misfit between the function $y(x, \mathbf{w})$, for any given value of \mathbf{w} , and the training set data points.
- Error Function: the sum of the squares of the errors between the predictions $y(x_n, \mathbf{w})$ for each data point x_n and the corresponding target values t_n .




$$E(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \{y(x_n, \mathbf{w}) - t_n\}^2$$


Polynomial Curve Fitting

- We can solve the curve fitting problem by choosing the value of w for which $E(w)$ is as small as possible.
- Since the error function is a quadratic function of the coefficients w , its derivatives with respect to the coefficients will be linear in the elements of w , and so the minimization of the error function has a unique solution, denoted by w^* ,
- The resulting polynomial is given by the function $y(x, w^*)$.
- Choosing the order M of the polynomial → model selection.


0th Order Polynomial




1st Order Polynomial



3rd Order Polynomial

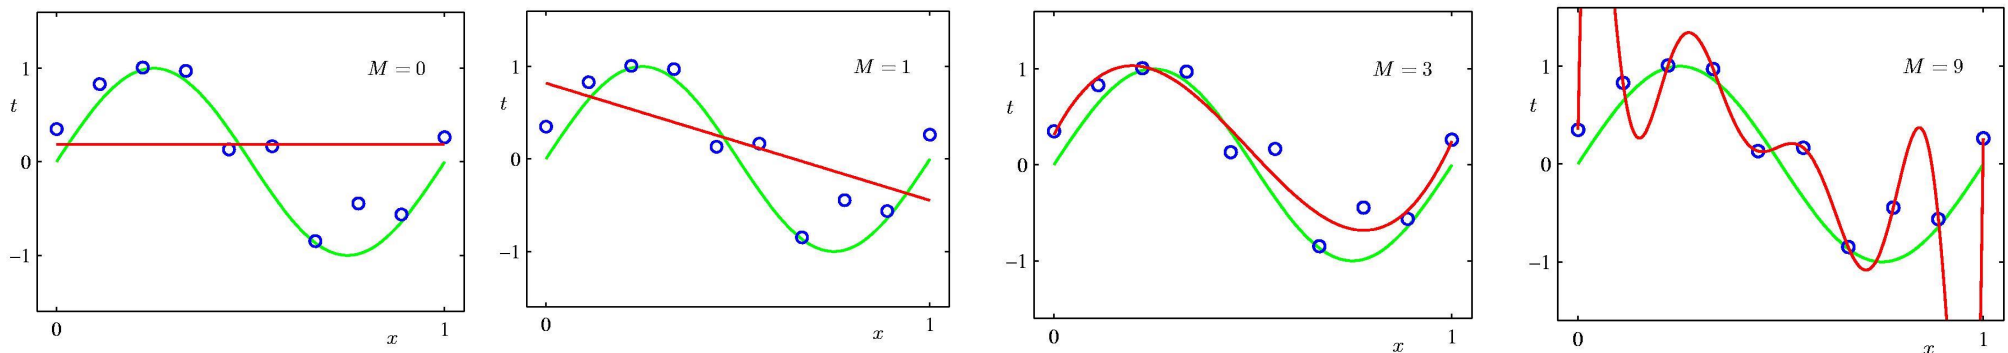


9th Order Polynomial




Polynomial Curve Fitting

- The 0th order ($M=0$) and first order ($M=1$) polynomials give rather poor fits to the data and consequently rather poor representations of the function $\sin(2\pi x)$.
- The third order ($M=3$) polynomial seems to give the best fit to the function $\sin(2\pi x)$ of the examples.
- When we go to a much higher order polynomial ($M=9$), we obtain an excellent fit to the training data.
 - In fact, the polynomial passes exactly through each data point and $E(w^*) = 0$.



Polynomial Curve Fitting

- We obtain an excellent fit to the training data with 9th order.
- However, the fitted curve oscillates wildly and gives a very poor representation of the function $\sin(2\pi x)$.




- This behaviour is known as **over-fitting**

Polynomial Curve Fitting

Over-fitting

- We can then evaluate the residual value of $E(\mathbf{w}^*)$ for the training data, and we can also evaluate $E(\mathbf{w}^*)$ for the test data set.
- Root-Mean-Square (RMS) Error: $E_{\text{RMS}} = \sqrt{2E(\mathbf{w}^*)/N}$
 - in which the division by N allows us to compare different sizes of data sets, and the square root ensures that E_{RMS} is measured on the same scale as the target variable t .



Polynomial Curve Fitting



Polynomial Coefficients

- Magnitude of coefficients increases dramatically as order of polynomial increases.
- Large positive and negative values so that the corresponding polynomial function matches each of the data points exactly, but between data points the function exhibits the large oscillations → over-fitting

	$M = 0$	$M = 1$	$M = 3$	$M = 9$
w_0^*	0.19	0.82	0.31	0.35
w_1^*		-1.27	7.99	232.37
w_2^*			-25.43	-5321.83
w_3^*			17.37	48568.31
w_4^*				-231639.30
w_5^*				640042.26
w_6^*				-1061800.52
w_7^*				1042400.18
w_8^*				-557682.99
w_9^*				125201.43

Polynomial Curve Fitting

- Increasing the size of the data set reduces the over-fitting problem.
- 9th Order Polynomial.



Polynomial Curve Fitting

regularization

- We may wish to use relatively complex and flexible models with data sets of limited size.
- The over-fitting phenomenon can be controlled with **regularization**, which involves adding a penalty term to the error function.
- **Regularization:** Penalize large coefficient values

$$\tilde{E}(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \{y(x_n, \mathbf{w}) - t_n\}^2 + \frac{\lambda}{2} \|\mathbf{w}\|^2$$

$$\text{where } \|\mathbf{w}\|^2 \equiv \mathbf{w}^T \mathbf{w} = w_0^2 + w_1^2 + \dots + w_M^2$$


- the coefficient λ governs the relative importance of the regularization term compared with the sum-of-squares error term.

Polynomial Curve Fitting


regularization

- Plots of $M = 9$ polynomials fitted to the data set using the regularized error function

no regularization ($\lambda=0$)



too much regularization




Polynomial Curve Fitting

regularization

- Polynomial Coefficients

	$\ln \lambda = -\infty$	$\ln \lambda = -18$	$\ln \lambda = 0$
w_0^*	0.35	0.35	0.13
w_1^*	232.37	4.74	-0.05
w_2^*	-5321.83	-0.77	-0.06
w_3^*	48568.31	-31.97	-0.05
w_4^*	-231639.30	-3.89	-0.03
w_5^*	640042.26	55.28	-0.02
w_6^*	-1061800.52	41.32	-0.01
w_7^*	1042400.18	-45.95	-0.00
w_8^*	-557682.99	-91.53	0.00
w_9^*	125201.43	72.68	0.01



Graph of the root-mean-square error versus $\ln \lambda$ for the $M=9$ polynomial.